

Abstract

This is the report for Exercise 3. In this report we showcase a deployment procedure for a Cloud system, as described in the [GitHub repository](#).

1 Introduction

This report describes the deployment of a cloud-based file storage system based on Nextcloud. The system allows users to authenticate and manage their own files, including uploading, downloading and deleting files. The system is deployed locally using Docker and Docker Compose, with Nextcloud providing the file storage and web interface and MariaDB providing the application database. Persistent Docker volumes are used to ensure that application data and files survive container restarts.

2 Architecture

The deployment consists of two Docker containers. The first container runs Nextcloud and provides the web interface and file-management functionality. The second container runs MariaDB and stores Nextcloud's application data.

The containers communicate through the Docker Compose network. The MariaDB container is not exposed directly to the host, while the Nextcloud container exposes its HTTP service through port 7489.

Two persistent Docker volumes are used. The `db_data` volume stores the MariaDB database, while `nextcloud_data` stores the Nextcloud installation and user files. Consequently, removing and recreating the containers does not remove the persistent data.

Figure 1 illustrates the main components of the deployment and their relationships:




Figure 1: Architecture of the Nextcloud deployment.

3 Deployment

The system is deployed using Docker Compose. The complete deployment configuration is provided in the repository in the `docker-compose.yml` file.

The deployment uses the following images:

- `nextcloud:stable`, providing the cloud storage application;
- `mariadb:11.8.9`, providing the relational database.

The database is initialized using environment variables containing the database name and credentials. Similarly, the initial Nextcloud administrator account is configured through environment variables.

Sensitive values are not committed to the repository. Instead, the repository contains a `.env.example` file which documents the required configuration variables. The actual `.env` file is excluded from version control.

The system can be deployed by first creating the environment configuration and then starting the containers:

```
cp .env.example .env
docker compose up -d
```

The running containers can be inspected with:

```
docker compose ps
```

The Nextcloud web interface is available at `http://localhost:7489`, or at the IP of the local device it's hosted on.

4 User and File Management

Nextcloud provides the authentication and authorization mechanisms required by the exercise. Users can authenticate through the web interface and each user has a separate file space.

The administrator account can create and manage users and therefore provides the administrative functionality required by the assignment. Regular users can upload, download and delete files belonging to their own storage space.

No custom implementation of authentication or file management was required, since these functionalities are provided by Nextcloud.

5 Security

5.1 Authentication and Authorization

Authentication and authorization are delegated to Nextcloud. Users must authenticate before accessing their files, while administrative operations are restricted to administrator accounts.

5.2 Credential Management and permissions

Database credentials and the initial administrator password are provided through environment variables rather than being hard-coded in the Docker Compose configuration. The actual `.env` file is not committed to the repository. Its permissions could also be restricted at the operating system level to prevent unauthorized access. Ideally, database and administrator credentials should be managed using a dedicated secrets mechanism, such as [Docker Secrets](#).

As an additional security measure, the deployment could also be implemented [using Docker in rootless mode](#). This would allow the containers to run without requiring root privileges on the host, reducing the privileges available to container processes and limiting the potential impact of a compromised container.

5.3 Network Security and HTTPS

The MariaDB service is only reachable from the Docker Compose network and is not directly exposed through a host port. This reduces the external attack surface of the deployment by preventing direct connections to the database from outside the container network.

The local deployment currently uses HTTP because it is intended for easy testing. For a production deployment, HTTPS should be provided through a reverse proxy such as Caddy, Traefik, or Nginx. This would ensure that communication between clients and the Nextcloud instance is encrypted.

5.4 Private Network Access

Another possible deployment model is to use a private VPN such as Tailscale **to expose the service only to authorized devices**. In this scenario, the Nextcloud instance would not need to be directly exposed to the public Internet, reducing its external attack surface. Access to the service could instead be restricted to devices connected to the private network, while Nextcloud would continue to provide its own application-level authentication and authorization.

This approach would be particularly suitable for a private deployment where the storage service is intended to be accessed only by a controlled set of users and devices. For a publicly accessible production service, HTTPS and an appropriate reverse-proxy setup would still be preferable.

5.5 Additional Security Features

Additional Nextcloud applications can be used to extend the security and functionality of the platform. For example, the **End-to-End Encryption application** can provide additional protection for selected files.

6 Performance Considerations

The performance of the system depends primarily on the available CPU, memory, storage and network bandwidth. File operations involving small files are likely to be more affected by request and metadata overhead, while large files are more dependent on storage and network throughput.

For a larger deployment, the performance could be evaluated by generating concurrent upload and download requests and measuring throughput, latency and resource utilization. The same approach could be used to compare workloads involving small, medium and large files.

The current deployment is intended as a local demonstration and does not attempt to reproduce a production-scale workload.

7 Scalability

The current deployment uses a single Nextcloud container and a single MariaDB container, which is sufficient for a local deployment but would not be suitable for a large number of users.

In a production environment, the system could be deployed using an Infrastructure as a Service (IaaS) provider such as AWS, Microsoft Azure or Google Cloud. Virtual machines or cloud container services could be used to run multiple instances of the Nextcloud application, with a load balancer distributing incoming requests between them. This would allow the application layer to scale horizontally as the workload grows.

The database could be moved to a managed or highly available database service, while user files could be stored using scalable object storage. Separating the application, database and file-storage layers would allow each component to be scaled independently.

Caching could also be introduced to reduce repeated database and application requests. Monitoring resource utilization and application performance would allow resources to be increased or reduced according to the actual workload.

8 Cost Efficiency

The proposed solution is based entirely on open-source software, such as Nextcloud, MariaDB and Docker, resulting in no software licensing costs.

The use of lightweight containers also helps keep resource usage low, while monitoring can be configured to provide the necessary information without introducing significant overhead. The local Docker deployment is therefore particularly suitable for development and testing, as it does not require any cloud infrastructure costs.

For a production deployment, cloud resources could be optimized according to the workload. For example, auto-scaling groups could dynamically adjust the number of application instances, while spot instances could be used for workloads that can tolerate interruptions. These approaches can reduce infrastructure costs while still allowing the system to handle periods of increased demand.